

IndiSpeak: Multilingual Language Detection using Feature-based Audio Classification

MSDA CAPSTONE
PROJECT FOR
SUMMER 2023


BY SONAL GANVIR

MENTOR:

DR.DVIJESH SHASTRI

Language recognition is a critical problem domain with immense practical applications.

- India is a country known for its rich linguistic diversity, with 22 languages officially recognized by the Indian Constitution [1].
- Language identification is essential for various applications, including:
 - **Customer Support:** Providing personalized assistance in the customer's native language for better user experience.
 - **Voice Assistants:** Enabling voice assistants to comprehend and respond in multiple Indian languages.
 - **Language Research:** Facilitating linguistic studies and preserving indigenous languages.



The dataset comprises a vast collection of audio samples representing 10 diverse Indian languages.

Massive Collection of Audio Samples

- 250k + voice samples
- mp3 format

Standardized Duration

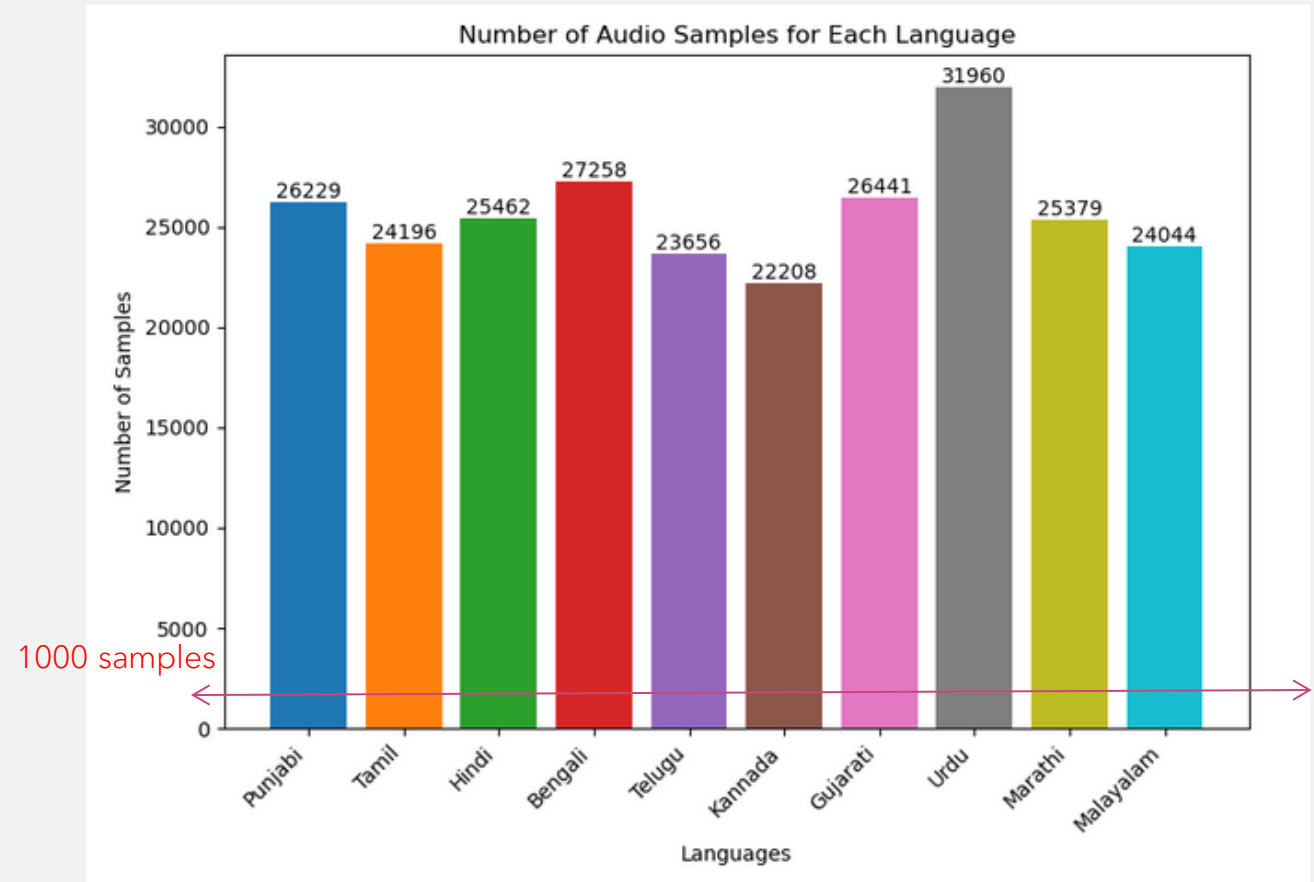
- Each voice sample ~ 5 sec long

Source of Dataset

- Extracted from YouTube videos

Supported Languages

- 10 Indian Languages : Punjabi, Tamil, Hindi, Bengali, Telugu, Kannada, Gujarati, Urdu, Marathi, Malayalam



Tools and Technologies Used

Language

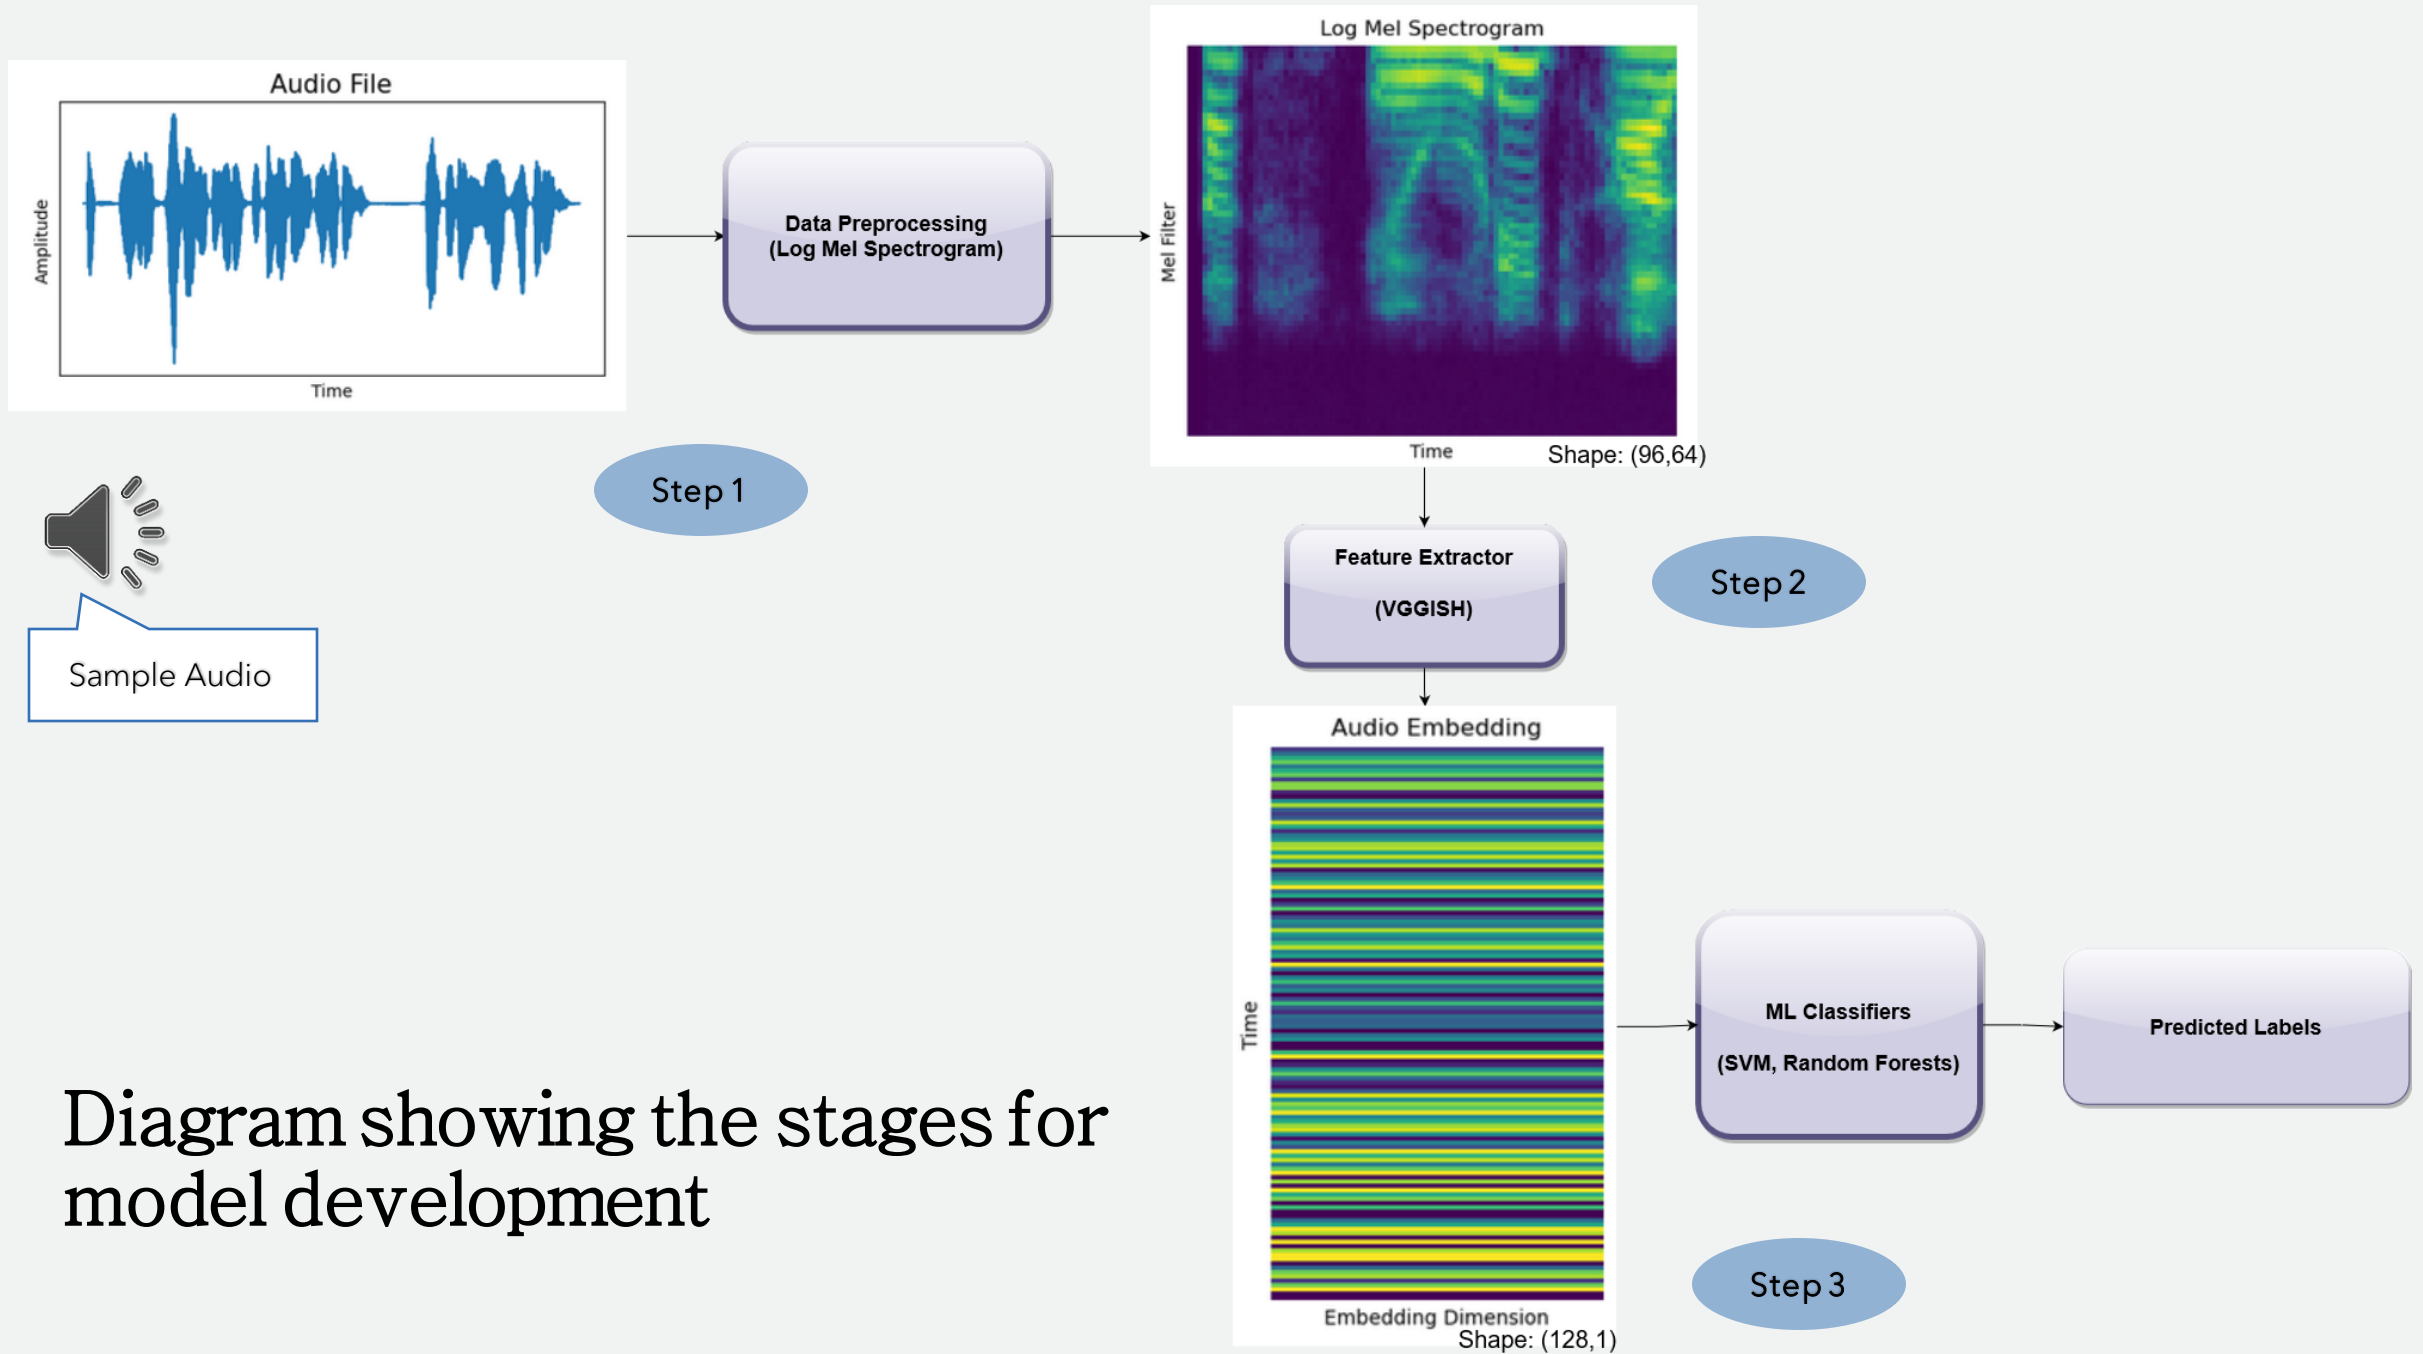


IDE

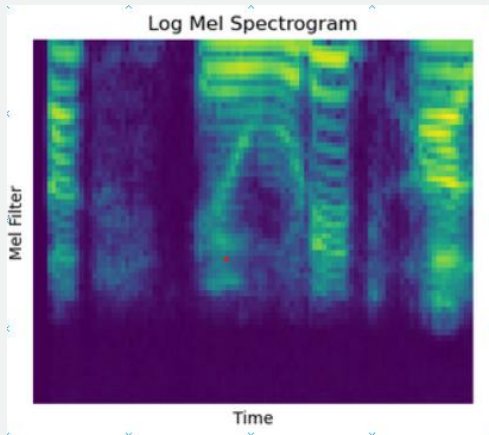


Packages






Converting the audio classification problem to an image classification task enhances language recognition accuracy.

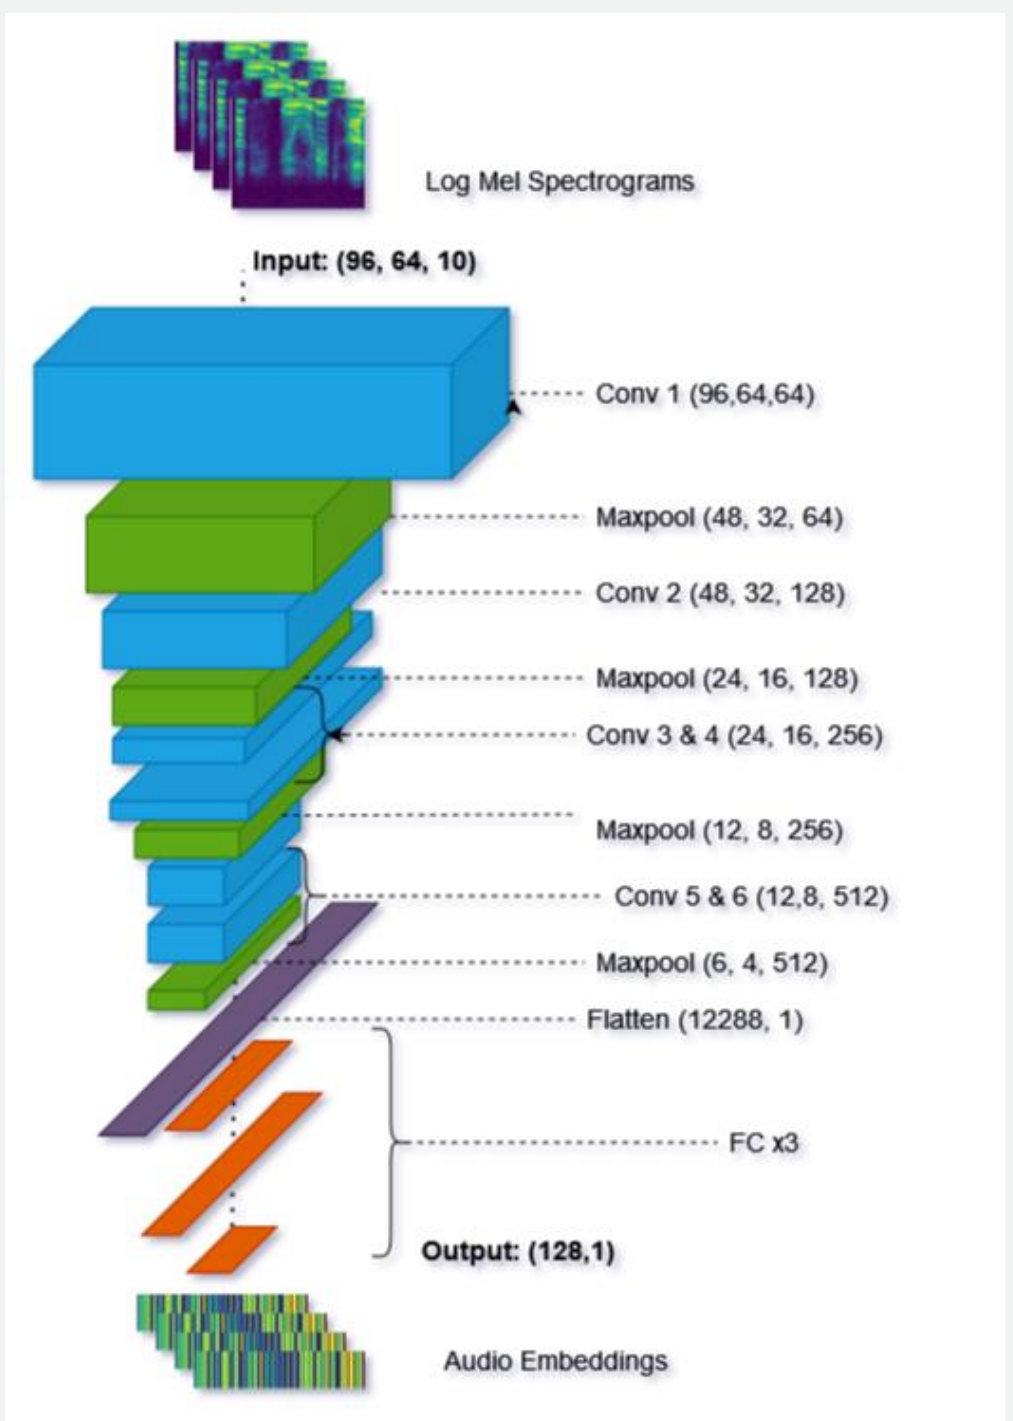


Mel Spectrogram: A mel spectrogram is a visual representation of the spectrum of frequencies in an audio signal, with frequencies on the y-axis and time on the x-axis.



VGGish Model – A Powerful Feature Extractor

- **VGGish Model:** VGGish is a pre-trained deep learning model developed by Google's Magenta team.
- It was originally designed for audio feature extraction, particularly for audio classification tasks.
- **Architecture:** VGGish is based on the VGG-16 architecture, modified for audio processing.
- The model consists of 6 convolutional layers and 4 max-pooling layers.
- The final layer is a fully connected layer with 128 units, followed by a ReLU activation.




Two powerful Machine Learning (ML) classifiers are employed for language identification.

1. Support Vector Machine (SVM) Classifier:

- **Working:** SVM is a supervised learning algorithm that finds an optimal hyperplane to separate data points into different classes, maximizing the margin between classes.
- **Usage:** SVM is applied to classify audio embeddings into Indian languages, effectively handling high-dimensional data and multiclass classification tasks.
- **Advantages:** SVM is robust against overfitting, efficient in handling large feature spaces, and performs well with limited training data.
- **Default Parameters:**

Trained Model Parameters:

C: 1.0 kernel: rbf degree: 3 gamma: scale

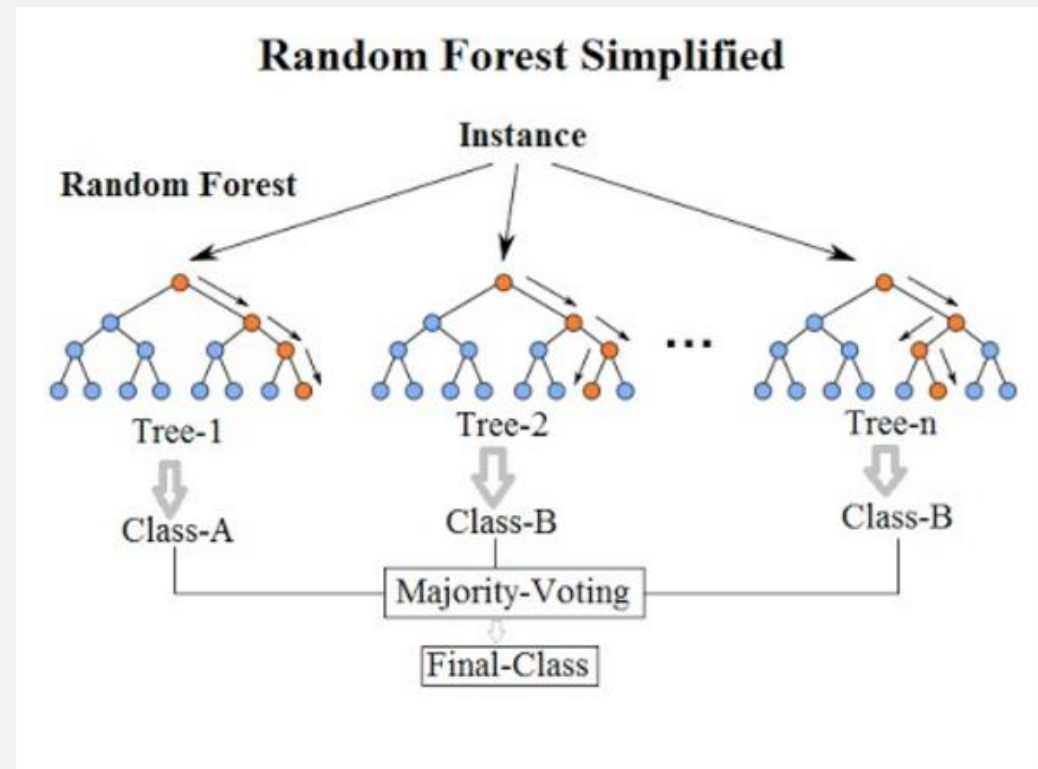


Two powerful Machine Learning (ML) classifiers are employed for language identification.

2. Random Forest (RF) Classifier:

- **Working:** Random Forest (RF) is an ensemble learning method that combines multiple decision trees, trained on random subsets of data and features, using majority voting for final predictions.
- **Usage:** RF is utilized to classify audio embeddings into Indian languages, benefiting from its ensemble nature that provides robustness and reduces overfitting risk.
- **Advantages:** RF effectively handles high-dimensional data, captures complex relationships, and performs well with imbalanced datasets.
- **Default Parameters:**

```
Trained Model Parameters:  
n_estimators: 100 max_depth: None min_samples_split: 2
```



Model Training and Evaluation

1. Training Process:

1. Samples Range: We train both classifiers for sample sizes ranging from 1,000 to 10,000 samples.
2. Data Split: 80% of the samples are used for training, ensuring the classifiers learn from a substantial portion of the data.
3. Class Balancing: To prevent bias towards dominant classes, we ensure class balancing during training.
4. Hyperparameter Tuning: We used GridSearchCV with cross-validation (cv=10) to find the best combination of hyperparameters for the model..

2. Evaluation Process:

1. Testing Set: The remaining 20% of samples serve as the testing set, ensuring an unbiased evaluation.
2. Performance Metrics: The classifiers are evaluated using accuracy, precision, recall, and F1-score to measure their effectiveness in language detection.
3. Confusion Matrix: Confusion matrices are generated to visualize classifier predictions and identify misclassifications.

3. Time and Space Complexity Analysis:

4. Kaggle Compute Instance: NVIDIA T4 (X2) with 30 GB RAM

1. Execution Time: We measure the time taken by each classifier to train and make predictions.
2. Memory Usage: We monitor the memory consumption during the execution of the classifiers.

In [7]:

```
!python main.py --model='random_forest' --run_name='run1' --total_samples=1000 --samples_per_
folder=100
```

Run Name: run0

ML Model: svm

GPU is available

Pre-trained model and files don't exist, start training

Training started.

Generating Train Embeddings: 1batch [03:38, 218.27s/batch]

Train Embeddings shape: (800, 640)

Train Labels shape: (800,)

Fitting the model...

Model fitting completed.

Label encoder saved.

Generating Test Embeddings: 1batch [00:49, 49.94s/batch]

Test Embeddings shape: (200, 640)

Accuracy: 0.73

Precision: 0.7798667172370483

Recall: 0.73

F1 Score: 0.7217700882906867

Sample Output

Confusion Matrix:

```
[[21  1  0  0  0  0  0  0  1  0]
 [ 0  9  0  0  0  0  1  0  0  1]
 [ 0  0 21  0  1  0  0  0  0  1]
 [ 0  0  0 15  1  1  0  0  0  1]
 [ 1  0  0  0 20  0  0  0  0  0]
 [ 1  0  0  0  0 17  0  3  0  1]
 [ 1 17  0  0  0  0  4  0  0  0]
 [ 1  0  0  0  0  1  0 15  0  1]
 [ 4  0  0  0  0  0  0  0 13  3]
 [ 2  0  0  1  0  8  0  0  0 11]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.91	0.78	23
1	0.33	0.82	0.47	11
2	1.00	0.91	0.95	23
3	0.94	0.83	0.88	18
4	0.91	0.95	0.93	21
5	0.63	0.77	0.69	22
6	0.80	0.18	0.30	22
7	0.83	0.83	0.83	18
8	0.93	0.65	0.76	20
9	0.58	0.50	0.54	22
accuracy			0.73	200
macro avg	0.76	0.74	0.71	200
weighted avg	0.78	0.73	0.72	200

Execution Time: 271.53480553627014 seconds

Memory Usage Before Execution: 1697.5390625 MB

RAM Usage Before Execution: 2158.2734375 MB

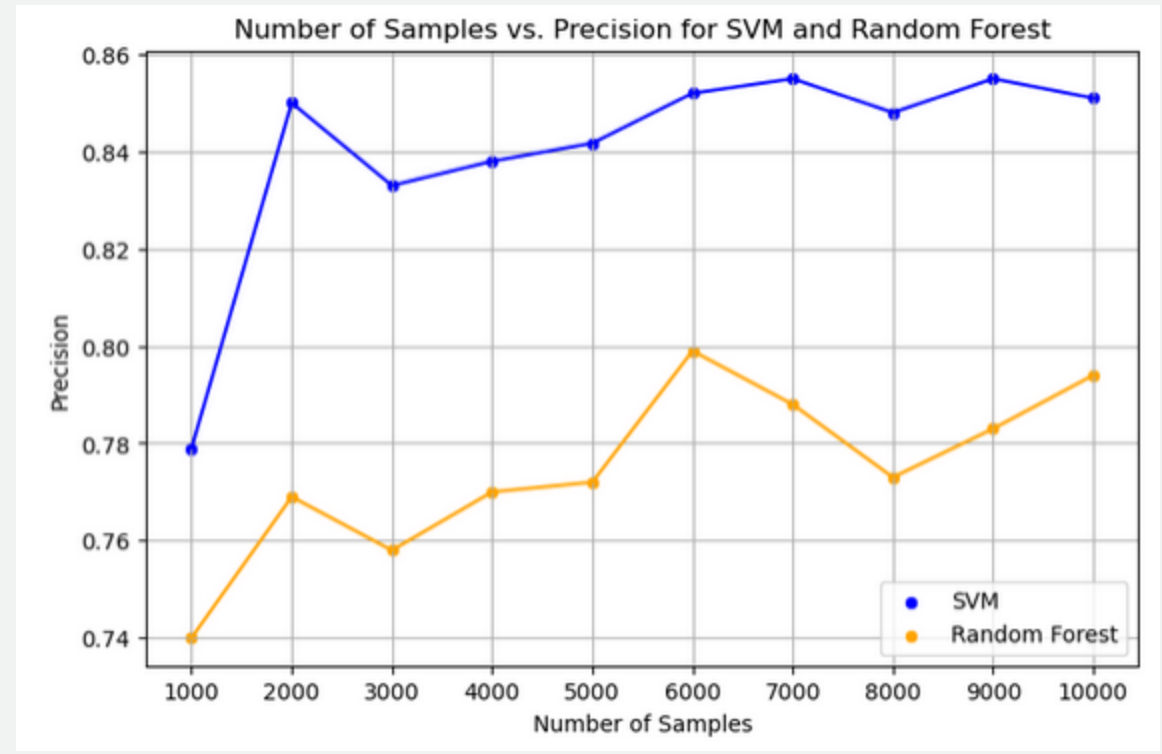

Memory Usage After Execution: 4182.875 MB

RAM Usage After Execution: 3746.5625 MB


Garbage collected: 30078 objects

Sample Output


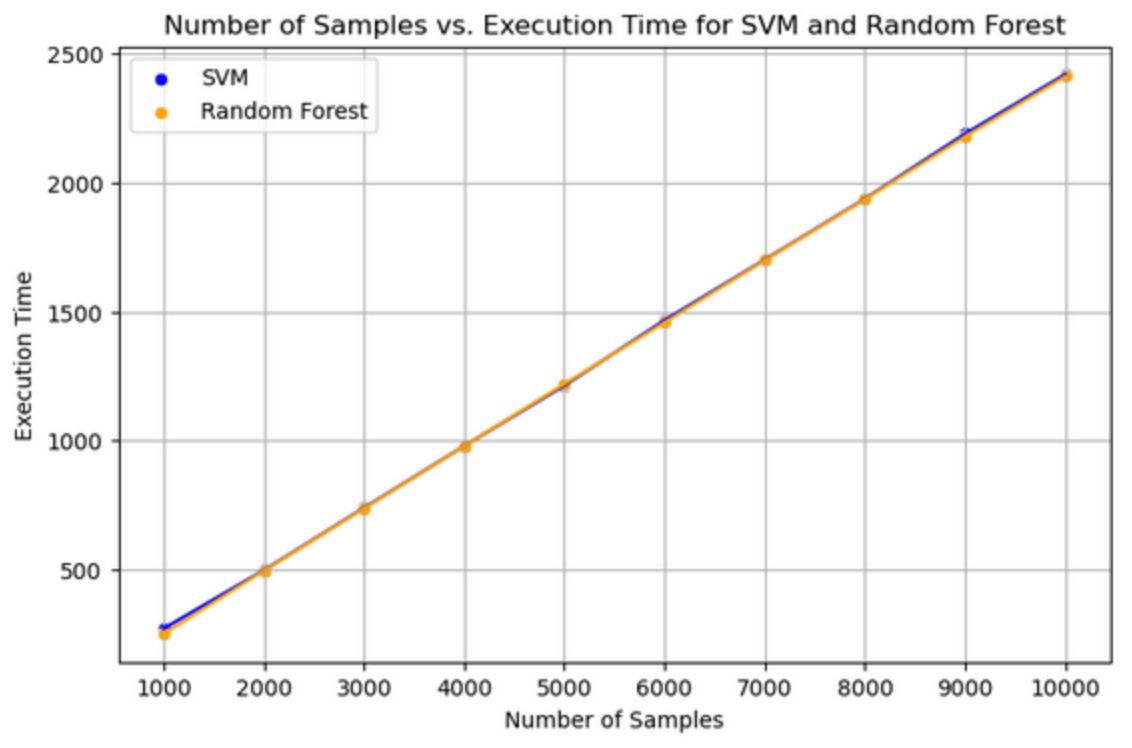
Performance Metrics (Accuracy + Precision)



Performance Metrics (Recall + F1-Score)



Performance Metrics (Execution Time + RAM Usage)



Hyperparameter Tuning

```
param_grid_svm = {  
    'C': [0.1, 1, 10],  
    'kernel': ['linear', 'rbf'],  
    'degree': [3, 5],  
    'gamma': ['scale', 'auto']  
}
```

```
Run Name: gsearch_svm  
ML Model: svm  
GPU is available  
Pre-trained model and files don't exist, start training  
Performing hyperparameter tuning...  
Generating Train Embeddings: 1batch [22:48, 1368.13s/batch]  
Train Embeddings shape: (5600, 640)  
Train Labels shape: (5600,)  
Hyperparameter tuning completed.  
Best Hyperparameters:  
C: 10  
degree: 3  
gamma: scale  
kernel: rbf  
Best accuracy Score: 0.8551785714285713
```

```
param_grid_rf = {  
    'n_estimators': [100, 200, 300],  
    'max_depth': [None, 10, 20],  
    'min_samples_split': [2, 5],  
    'min_samples_leaf': [1, 2]  
}
```

```
Run Name: gsearch_rf  
ML Model: random_forest  
GPU is available  
Pre-trained model and files don't exist, start training  
Performing hyperparameter tuning...  
Generating Train Embeddings: 1batch [20:13, 1213.60s/batch]  
Train Embeddings shape: (4800, 640)  
Train Labels shape: (4800,)  
Hyperparameter tuning completed.  
Best Hyperparameters:  
max_depth: None  
min_samples_leaf: 1  
min_samples_split: 5  
n_estimators: 300  
Best accuracy Score: 0.8006249999999999
```

ML Training Summary

- Total Models Built: 10 models for 2 ML Algorithm i.e. 20 + 2 GridSearchCV Models
- Minimum Execution Time = 4 min
- Maximum Execution Time = 40 min
- Average Execution time for a regular ML model = 22 min

- Average Execution Time for hyperparameter tuning = 1 h 48 min
- Total Combinations tried for SVM = 24
- Total Combinations tried for RF = 36

Conclusion – Best Performing Model

```
Training started.  
Generating Train Embeddings: 1batch [22:40, 1360.33s/batch]  
Train Embeddings shape: (5600, 640)  
Train Labels shape: (5600,)  
Fitting the model...  
Trained Model Parameters:  
C: 10 kernel: rbf degree: 3 gamma: scale  
Label encoder saved.  
Generating Test Embeddings: 1batch [05:40, 340.30s/batch]  
Test Embeddings shape: (1400, 640)  
Accuracy: 0.8564285714285714  
Precision: 0.8557778379729106  
Recall: 0.8564285714285714  
F1 Score: 0.855740447796386
```

- **Best Model:** *gsearch_svm* (SVM Classifier)
- **Accuracy:** 0.85
- **Precision:** 0.85
- **Recall:** 0.85
- **F1 Score:** 0.85
- **Note:** The selected model, 'gsearch_svm', demonstrates high accuracy and a good balance between precision and recall, making it a reliable and effective choice for language detection in real-world applications.

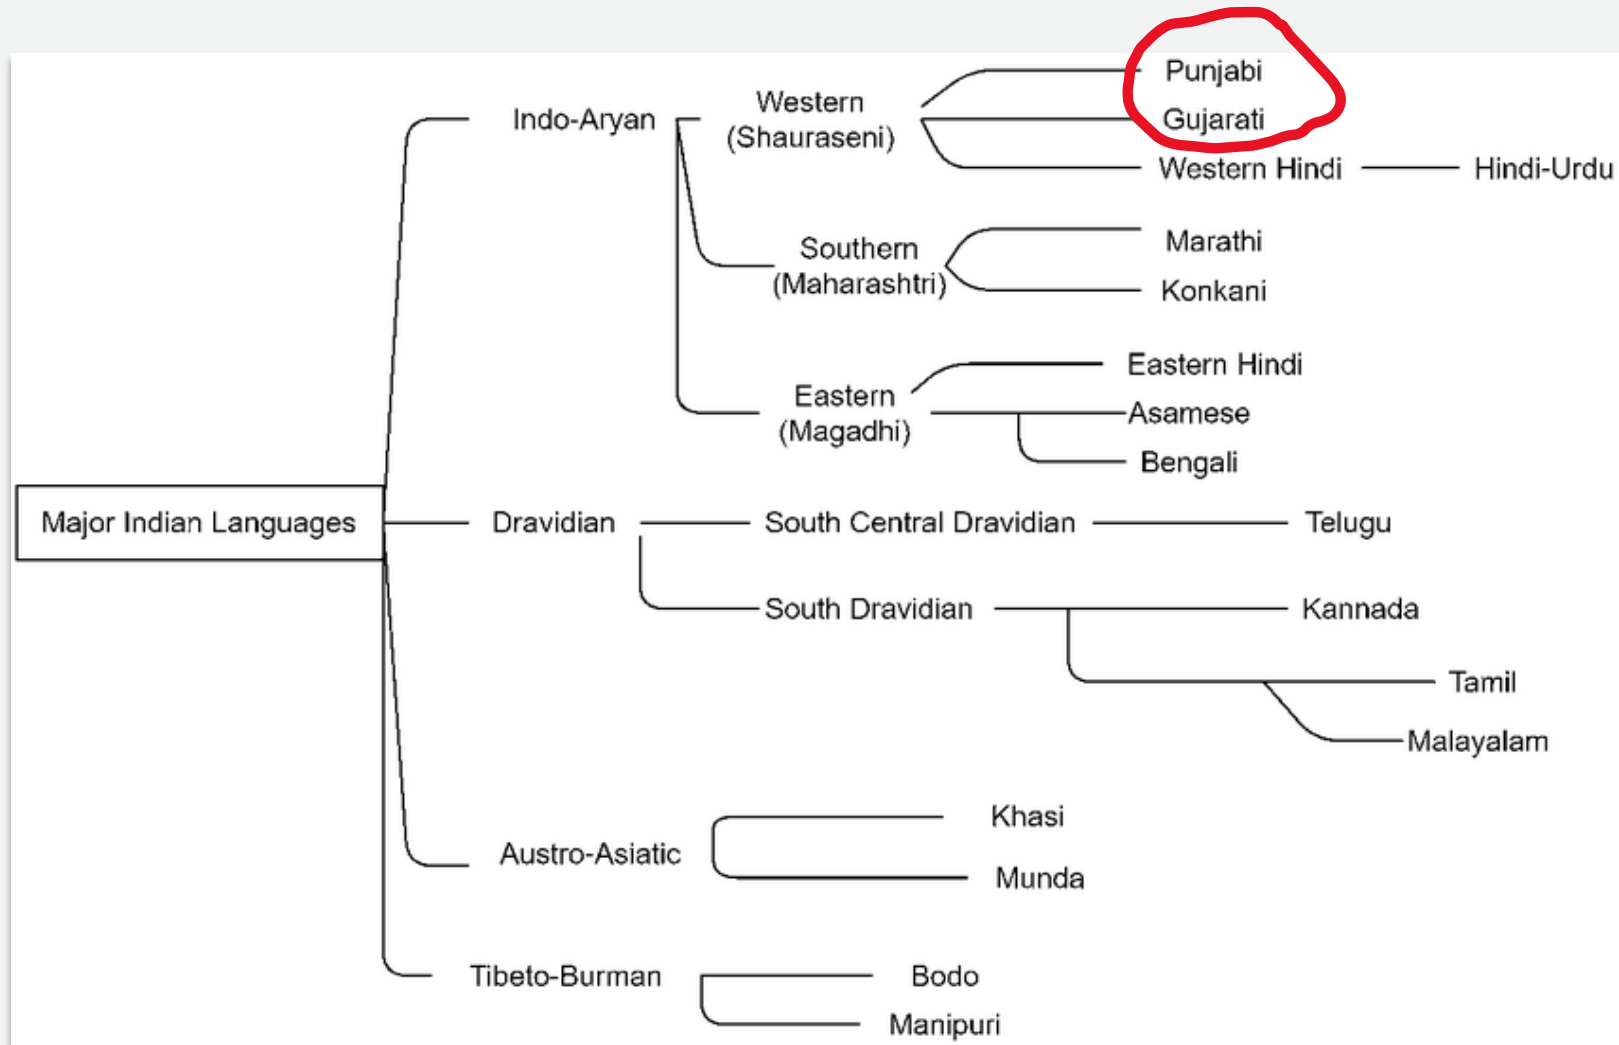
Performance Reports for *gsearch_svm*



Confusion Matrix

	precision	recall	f1-score	support
0	0.88	0.95	0.91	136
1	0.48	0.46	0.47	142
2	0.93	0.99	0.96	127
3	0.97	0.93	0.95	138
4	0.98	0.97	0.98	154
5	0.95	0.95	0.95	147
6	0.49	0.50	0.50	140
7	0.98	0.95	0.96	139
8	0.95	0.98	0.97	125
9	0.94	0.90	0.92	152

Classification Report



Tree diagram to illustrate the language closeness of major Indian languages

Acknowledgements

Thank You!


Ms. Angel Griffin (MSDA Associate Director)

Dr. Ermelinda DeLaVina (MSDA Director)

Dr. Dvijesh Shastri (Capstone Mentor)

All respectful MSDA Faculties!

Demo



Github Repo: <https://github.com/sonalغان/IndiSpeak>